Anthropic Partners with AMD to Deploy 2 Gigawatts of Compute Infrastructure

Anthropic announced a strategic compute agreement with AMD to expand its global AI data center footprint. The multi-year deal secures up to 2 gigawatts of dedicated compute capacity for training and serving next-generation artificial intelligence models.

The infrastructure deployment relies on next-generation AMD Instinct MI450 Series GPUs. These accelerators feature ultra-high memory bandwidth and advanced inter-chip interconnects. Anthropic will integrate these clusters directly into its core research environment.

Scaling Infrastructure for Claude Models

Training state-of-the-art frontier models demands massive computing power. As model sizes and reasoning capabilities grow, hardware bottlenecks can slow down scientific progress. Anthropic will use the new 2-gigawatt capacity to scale its upcoming Claude model families.

The AMD Instinct MI450 architecture delivers significant performance upgrades for large-scale cluster deployments. The chips offer high floating-point operations per second (FLOPS) while optimizing energy efficiency per token. This power efficiency helps reduce the overall environmental footprint of large data center operations.

Besides model training, the new hardware will support real-time inference workloads. Enterprise clients require fast API responses and low latency when running agentic workflows. High-density GPU clusters allow Anthropic to maintain high throughput during peak demand periods.

Diversifying Hardware Supply Chains

This partnership marks a major step toward hardware diversification for Anthropic. Historically, leading AI labs relied on a single dominant chip supplier for high-performance training. Relying on multiple hardware vendors reduces supply chain risks and lowers hardware deployment delays.

Anthropic collaborated closely with AMD engineers to optimize the ROCm software stack. Open-source software stack integration ensures that Anthropic's custom training frameworks run natively on AMD GPUs with minimal overhead. Code optimization efforts have already demonstrated competitive training speeds across standard benchmarks.

> Securing multi-gigawatt compute commitments ensures long-term infrastructure stability for frontier AI development.

Data Center Architecture and Expansion Timeline

The physical deployment will span multiple energy-efficient data center facilities across North America. The facilities utilize liquid cooling technology to handle high power densities without thermal throttling. Liquid cooling also reduces overall cooling energy consumption compared to traditional air cooling.

The first 500-megawatt compute cluster will become operational later this year. Additional data center capacity will come online in planned phases through 2028. This phased expansion ensures continuous scaling as research teams introduce larger training runs.

The agreement underlines the growing demand for dedicated AI compute infrastructure worldwide. By securing large-scale GPU access early, Anthropic positions itself to sustain long-term research velocity and enterprise availability.

Contributor: Alessandro Linzi