

Claude Opus 5: Anthropic's Most Efficient Frontier Model Yet

Anthropic released **Claude Opus 5** on July 24, 2026, positioning it as the efficiency flagship of the Claude model family. The pitch is direct: near-Fable-5 performance at half the price — and in many agentic benchmarks, performance that doesn't just approach Fable 5 but matches or exceeds it.

Performance and Benchmarks

Opus 5 leads on coding and knowledge-work evaluations including Frontier-Bench and GDPval-AA, according to Anthropic's internal evaluations. On general agentic tasks and software engineering benchmarks, it frequently matches Claude Fable 5 — Anthropic's previous frontier model, which was specifically optimised for complex research and safety-critical use cases.

The key differentiator isn't raw capability in isolation — it's the combination of that capability with improved **reliability**. Opus 5 is built to complete multi-step agentic tasks without leaving stubs, abandoning mid-task, or requiring repeated re-prompting. It can verify its own work across long outputs and catch inconsistencies before a task is marked complete. This matters enormously in production agentic systems, where partial completions are often worse than no action at all.

Adjustable Reasoning Depth

One of the more technically interesting features of Opus 5 is its **adjustable computational effort system**. Users and API developers can set reasoning intensity across five levels: low, medium, high, xhigh, and max. Lower settings reduce latency and cost; higher settings engage deeper chains of verification and cross-checking before the model commits to an output.

This granularity is valuable in agentic pipelines where different subtasks warrant different levels of scrutiny. A file rename can run at low effort; a security-critical configuration change warrants max. Rather than forcing developers to choose between a fast, light model and a slow, heavy one, Opus 5 provides a continuous dial — with the same underlying model architecture adjusting its internal compute budget dynamically.

A single model that can operate at five levels of reasoning depth changes how you think about agentic pipeline design — one model, optimised per task, rather than a fleet of specialised models.

Calibrated Safety Without Over-Refusal

One friction point with Claude Fable 5 noted by developers was its tendency toward over-refusal — triggering safety classifiers on benign professional tasks and interrupting agentic workflows unexpectedly. Anthropic reports that Opus 5 intervenes approximately **85% less often** than Fable 5 on the same tasks, while maintaining equivalent safety standards on genuinely sensitive inputs.

This recalibration reflects an ongoing tension in frontier model development: safety interventions that are too aggressive don't just frustrate users, they actively prevent legitimate and beneficial use. Finding the right threshold — intervening on real risks while leaving professional workflows uninterrupted — is an alignment problem in itself, and Opus 5 represents Anthropic's current best answer.

Pricing and Availability

Claude Opus 5 is priced at **\$5 per million input tokens** and **\$25 per million output tokens** — identical to the previous Claude Opus 4.8, and approximately half the cost of Fable 5. The model supports a **1-million-token context window**, making it viable for large codebase analysis, long-form document review, and extended multi-turn agentic sessions.

It is now the default model for Claude Max subscribers and the highest-capability model available to Claude Pro users. API access is available through Anthropic's API directly, as well as through Amazon Bedrock, Google Cloud, and other major cloud platforms. For teams running agentic workflows at scale, the combination of frontier performance, adjustable reasoning, and halved cost relative to Fable 5 makes Opus 5 the most immediately practical model Anthropic has shipped.

Contributor: Alessandro Linzi