

Gemini Robotics 2: Google DeepMinds Vision for Whole-Body Robot Intelligence

Google DeepMind released **Gemini Robotics 2** on July 30, 2026 — a suite of three physical AI models that represent the most significant leap in embodied AI since the original Gemini Robotics launch. Where previous iterations focused primarily on robotic arm control and desktop-scale manipulation, the new release targets **full humanoid robots**: legs, torso, multi-fingered hands, and all.

Three Models for Three Roles

The Gemini Robotics 2 suite is structured around distinct functional roles, each optimised for a different position in a robot's computing stack.

Gemini Robotics 2 (VLA) is a Vision-Language-Action model that converts visual and linguistic input directly into motor control signals. It operates as a single learned policy over the entire humanoid body — meaning the same model that decides to pick up an object also controls the stepping required to reach it, the torso rotation to orient correctly, and the five-fingered grip to complete the action. DeepMind describes this as eliminating the traditional architecture of separate controllers for locomotion, manipulation, and grasping.

Gemini Robotics ER 2 (Embodied Reasoning) functions as the high-level cognitive layer — a “brain” that handles spatial, temporal, and physical reasoning. It plans multi-step tasks, communicates goals to human collaborators in natural language, and coordinates actions across teams of robots. ER 2 is particularly notable for enabling *multi-robot collaboration*: different robots running ER 2 share a semantic model of their environment and can delegate subtasks to one another without explicit human orchestration.

Gemini Robotics On-Device 2 is an optimised version of the VLA model designed to run on a robot's onboard computer without a continuous cloud connection. Low-latency local inference is critical for physical manipulation tasks, where a 100ms round-trip to a cloud server can translate into a dropped object or a misjudged footstep.

Whole-body control under a single learned policy is the architecture change that makes humanoid robots practical, not just impressive.

Cross-Platform Adaptability

One of the most commercially significant aspects of Gemini Robotics 2 is its platform-agnosticism. DeepMind reports that the models can be adapted to new robotic hardware — different joint configurations, sensor arrays, and actuator characteristics — in as little as a few hours. This lowers the barrier for hardware manufacturers to integrate Gemini-level intelligence without requiring DeepMind to train a separate model for each robot form factor.

The ability to adapt quickly across embodiments also matters for the research community. Labs running humanoid research platforms can now deploy a common intelligence layer across heterogeneous hardware fleets, enabling more systematic comparisons of physical robot capabilities.

Availability and Developer Access

Gemini Robotics ER 2 is currently accessible to developers via **Google AI Studio** and in private preview on the **Gemini Enterprise Agent Platform**. The VLA and On-Device models are available to early-access partners and a select group of trusted testers — primarily academic research labs and robotics companies that have signed partnership agreements with DeepMind.

DeepMind has not announced a timeline for general availability, though the structure of the release — developer access now, broader rollout later — mirrors the cadence used for previous Gemini model launches. The implication is that the research community gets first access while DeepMind refines safety protocols for wider deployment of physical AI in uncontrolled environments.

For anyone tracking where embodied AI is heading, Gemini Robotics 2 sets the benchmark: a unified intelligence layer that controls the full body of a humanoid robot, reasons about multi-step tasks in natural language, collaborates with other robots, and runs locally when the network isn't available.

Contributor: Alessandro Linzi