

DeepMind and the UK AI Security Institute: making safety more measurable

AI safety can sound like philosophy until it turns into something operational: shared access, shared methods, and shared measurements. That's what stood out in Google DeepMind's announcement about deepening its partnership with the UK AI Security Institute (AISI) through a new Memorandum of Understanding focused on foundational security and safety research.

The key shift is moving beyond "test the model" moments toward longer-running research collaboration — the kind of work that can accumulate into better evaluation tooling over time.

What the partnership includes

DeepMind frames the updated partnership around a few practical commitments:

- Sharing access to proprietary models, data, and ideas to accelerate research progress.
- Joint reports and publications to share findings with the broader research community.
- More collaborative security and safety research and ongoing technical discussions.

This is the boring-but-important infrastructure layer of safety: not just finding issues, but building the machinery to keep finding them.

Three areas they'll focus on

The announcement calls out three research directions that feel especially relevant as models become more capable:

- **Monitoring AI reasoning processes:** techniques for tracking a system's "thinking," often described as chain-of-thought monitoring, to better understand how answers are produced.
- **Social and emotional impacts:** work on "socioaffective misalignment," where a system can follow instructions but still behave in ways that don't align with human well-being.
- **Economic systems:** simulating real-world tasks, having experts score them, and using that to reason about longer-term labour market impacts.

None of this is a silver bullet, but it's a sign that frontier AI safety is increasingly treated like an engineering and measurement problem — something you can improve with better tools, not just better intentions.